

Reinforcement Learning for Network Management

MADRE

Guillaume Fraysse

Orange Research, Châtillon,
France

January 12th, 2026

Slides and code samples



slides and code at <https://github.com/gfraysse/MADRE>



Acknowledging the work of a whole team: current and past members

Current members from Orange team:

- Guillaume FRAYSSE
- David DELANDE
- Edgar FERNANDES
- Abhishek GUPTA
- Abhishek SAINI
- Shantanu VERMA

Past members:

- Manu CHAUHAN
- Francesca FOSSATI
- Imen GRIDA BEN YAHIA
- Xinqi LONG (intern)
- Yoichi MATSUO (NTT Network Service Systems Laboratories)
- Jose Manuel SANCHEZ VILCHEZ
- Jatinder SINGH
- Gabriel Gomes De SIQUEIRA (intern)



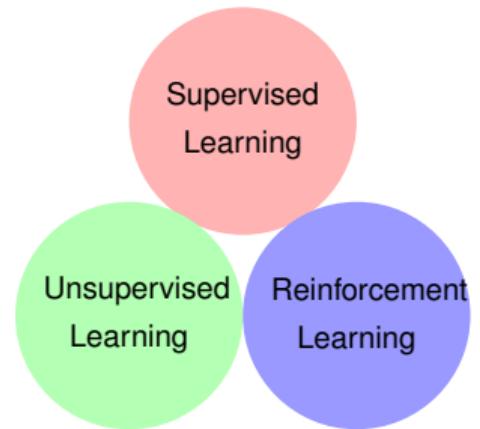
Related publications

Publications

- J. Singh, S. Verma, Y. Matsuo, F. Fossati and G. Fraysse, **Autoscaling Packet Core Network Functions with Deep Reinforcement Learning** NOMS 2023-2023 IEEE/IFIP Network Operations and Management Symposium, Miami, FL, USA, 2023, pp. 1-6, doi: 10.1109/NOMS56928.2023.10154312
- Y. Matsuo, J. Singh, S. Verma and G. Fraysse, **Integrating state prediction into the Deep Reinforcement Learning for the Autoscaling of Core Network Functions** NOMS 2023-2023 IEEE/IFIP Network Operations and Management Symposium, Miami, FL, USA, 2023, pp. 1-5, doi: 10.1109/NOMS56928.2023.10154301
- S. Verma and G. Fraysse, **Setting up a Reinforcement Learning pipeline for a Telco Core Network** 27th Conference on Innovation in Clouds Internet and Networks (ICIN) Tutorial, Paris, France, 2024
- X. Long and G. Fraysse, **Safe RL for Core Network autoscaling**, 2024 20th International Conference on Network and Service Management (CNSM), Prague, Czech Republic, 2024, pp. 1-7, doi: 10.23919/CNSM62983.2024.10814355.
- S. Verma, A. Gupta, J. Singh, J. M. Sanchez Vilchez, G. Fraysse, **Adaptive control policies for Core Network autoscaling with Meta-RL** 2025 21th International Conference on Network and Service Management (CNSM), Bologna, Italy, 2025.

Definitions: the 3 main paradigms of Machine Learning

- **Machine Learning (ML)**: A computer program is said to learn from **experience E** with respect to some **class of tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E. (Prof. Tom M. Mitchell [Mitchell, 1997])

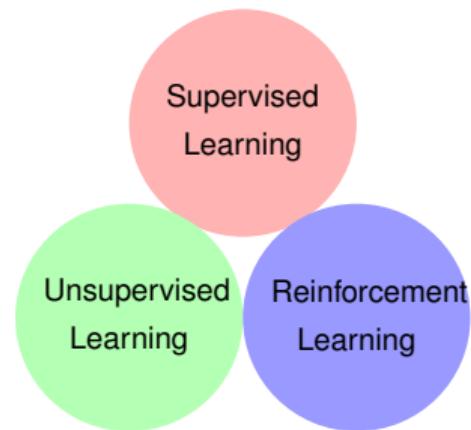


A Taxonomy of Machine Learning



Definitions: the 3 main paradigms of Machine Learning

- **Machine Learning (ML)**: A computer program is said to learn from **experience E** with respect to some **class of tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E. (Prof. Tom M. Mitchell [Mitchell, 1997])
 - **Supervised Learning**: Estimating some function $f : X \rightarrow Y$ given a set of labeled training examples (X_i, y_i) .



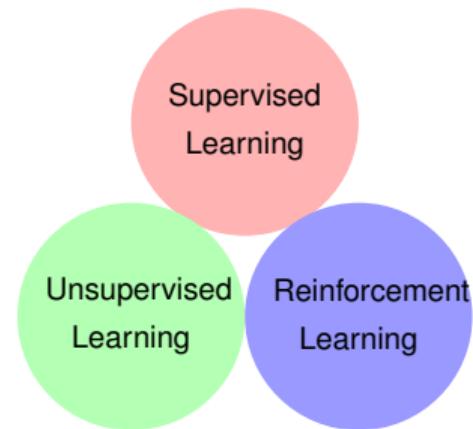
A Taxonomy of Machine Learning



Definitions: the 3 main paradigms of Machine Learning

- **Machine Learning (ML)**: A computer program is said to learn from **experience E** with respect to some **class of tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E. (Prof. Tom M. Mitchell [Mitchell, 1997])

- **Supervised Learning**: Estimating some function $f : X \rightarrow Y$ given a set of labeled training examples (X_i, y_i) .
- **Unsupervised Learning**: Harder to define formally. One common task is the **clustering**: construct a **grouping** of unlabeled data in **homogeneous** classes



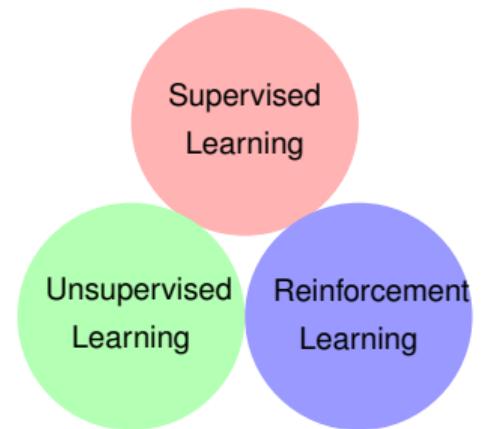
A Taxonomy of Machine Learning



Definitions: the 3 main paradigms of Machine Learning

- **Machine Learning (ML)**: A computer program is said to learn from **experience E** with respect to some **class of tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E. (Prof. Tom M. Mitchell [Mitchell, 1997])

- **Supervised Learning**: Estimating some function $f : X \rightarrow Y$ given a set of labeled training examples (X_i, y_i) .
- **Unsupervised Learning**: Harder to define formally. One common task is the **clustering**: construct a **grouping** of unlabeled data in **homogeneous** classes
- **RL**: learning what to do - how to map situations to actions - so as to maximize a numerical reward signal [Sutton and Barto, 2018]

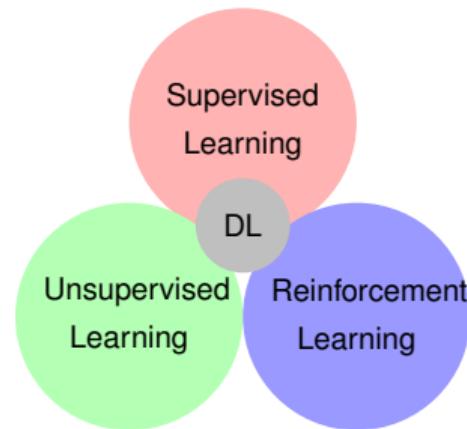


A Taxonomy of Machine Learning

Definitions: the 3 main paradigms of Machine Learning

- **Machine Learning (ML)**: A computer program is said to learn from **experience E** with respect to some **class of tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E. (Prof. Tom M. Mitchell [Mitchell, 1997])

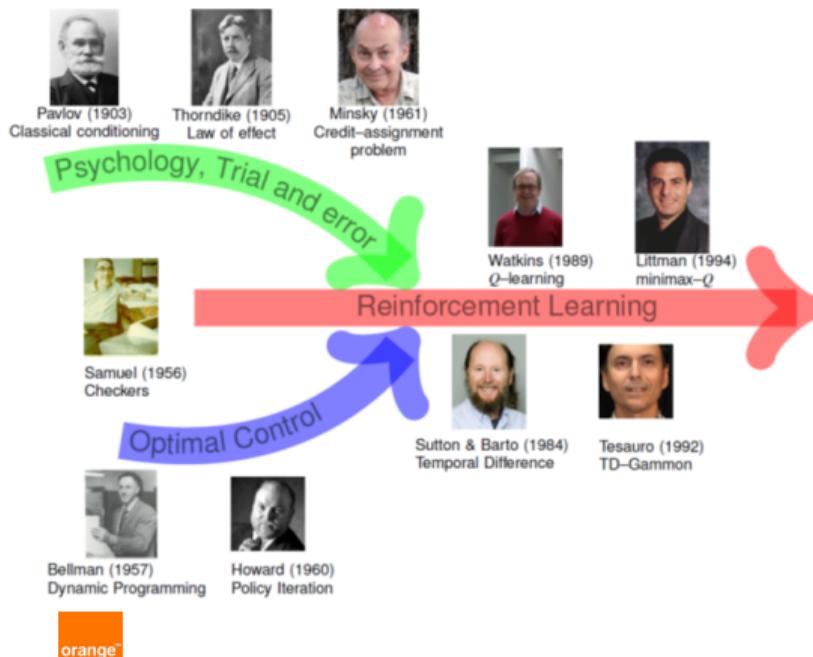
- **Supervised Learning**: Estimating some function $f : X \rightarrow Y$ given a set of labeled training examples (X_i, y_i) .
- **Unsupervised Learning**: Harder to define formally. One common task is the **clustering**: construct a **grouping** of unlabeled data in **homogeneous** classes
- **RL**: learning what to do - how to map situations to actions - so as to maximize a numerical reward signal [Sutton and Barto, 2018]



A Taxonomy of Machine Learning

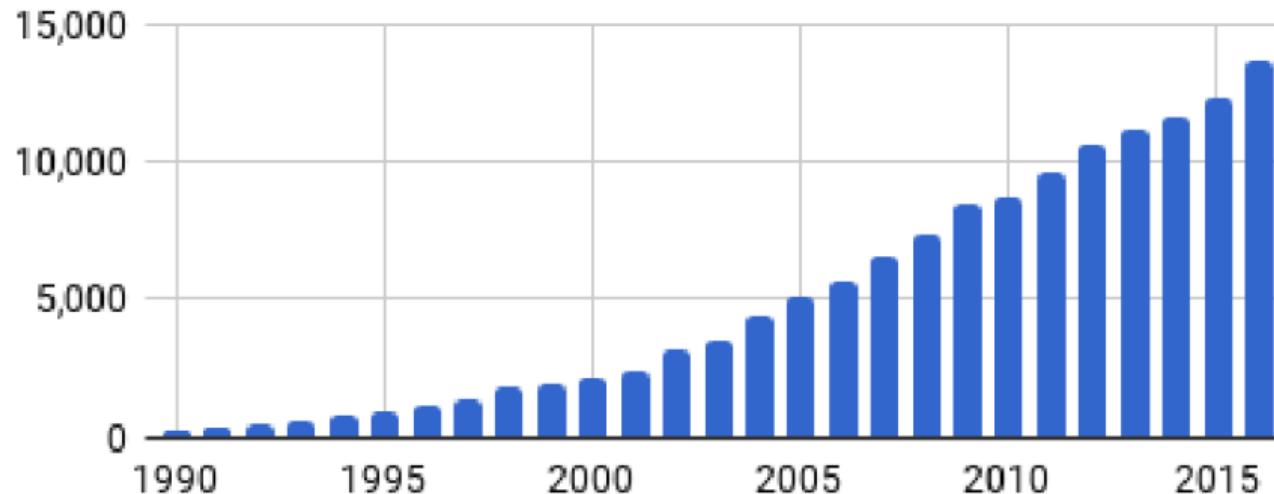
Deep Learning (DL): subset of ML methods based on Neural Networks

History of Reinforcement Learning



From Prof. Restelli
<https://sites.google.com>

Growing interest in RL



Evolution of the number of new RL-related publications per year, from [Henderson et al., 2018]



RL Application Areas 1/2

| Application Areas | Problem Statement | Environment |
|-------------------|--|--|
| Gaming | Chess, Pinball, Atari 2600 Games ([Mnih et al., 2013]), AlphaGo (2015) , AlphaStar ([Vinyals et al., 2019]), DOTA2 ([Berner et al., 2019]) | OpenAI/Farama Foundation (Atari, Classic Control), DeepMind Lab , ... |
| Cloud systems | Job scheduling, Congestion control, Database Optimization (Query optimization, Index structure), Resource Management and Autoscaling | Apache OpenWhisk FaaS ([Qiu et al., 2022]), Kubernetes cluster ([Qiu et al., 2020]), ... |

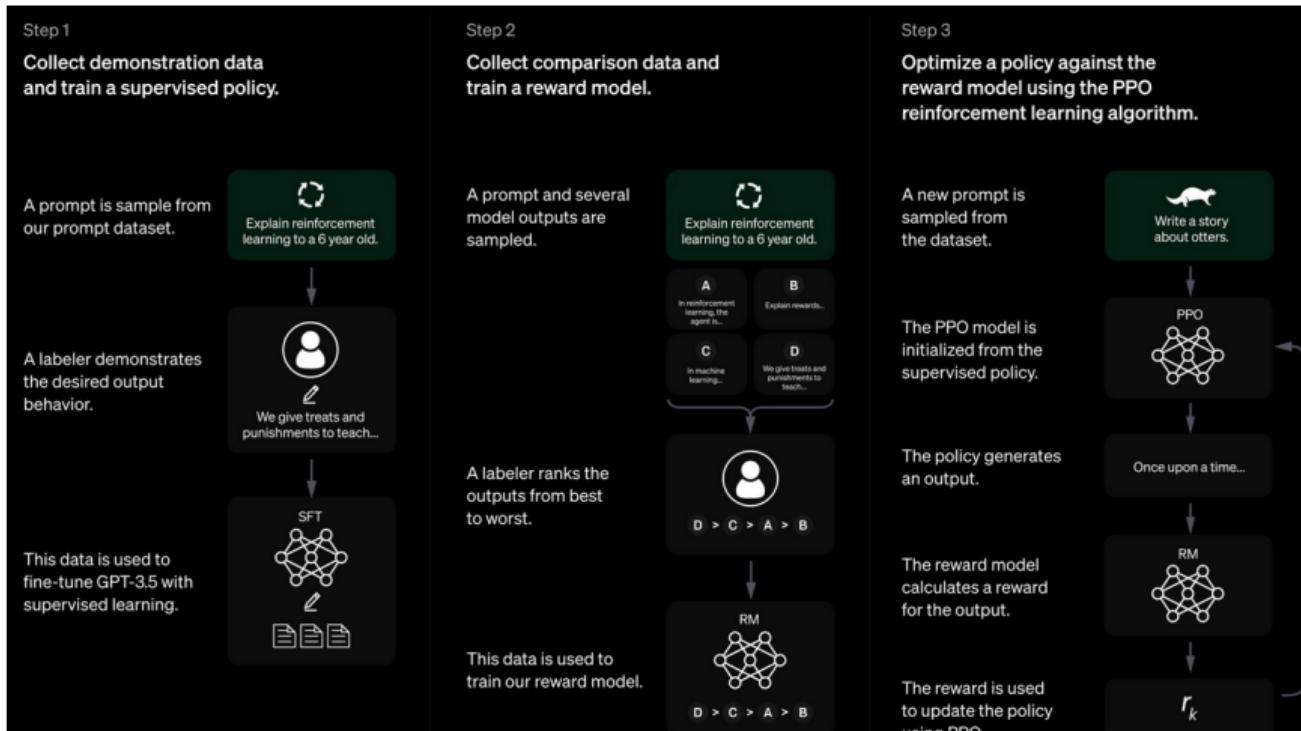


RL Application Areas 2/2

| Application Areas | Problem Statement | Environment |
|-------------------------|---|--|
| Autonomous driving cars | Trajectory optimization, dynamic pathing, Motion planning, route changing, decision position of parking | Carla, DeepTraffic by MIT |
| Robotics | Object Grasping, Obstacle Avoidance, Navigation & path planning, Simultaneous task execution | CopelliaSim, MuJoCo, PyRobot |
| Generative AI | Reinforcement Learning with Human Feedback (RLHF) for Large Language Models | Prompts, outputs and rankings by humans (Fine tuning of GPT models for ChatGPT, Llama2, ...) |

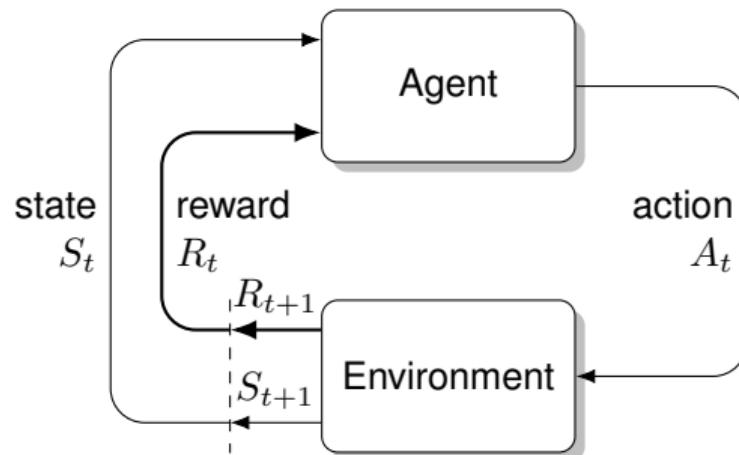


ChatGPT (from <https://openai.com/blog/chatgpt/>)



orange

Reinforcement Learning



Agent-environment interaction loop in a RL system, from [Sutton and Barto, 2018]



A basic example: moving in a maze



Maze (from Hugging Face tutorial)

| Element | Values set |
|-------------------|---|
| Action Space | direction the mouse goes to: \uparrow , \downarrow , \leftarrow or \Rightarrow |
| Observation Space | the cell the mouse is in |
| Reward | +0: Going to a state with no cheese in it. +1: Going to a state with a small cheese in it. +10: Going to the state with the big pile of cheese. -10: Going to the state with the poison. |

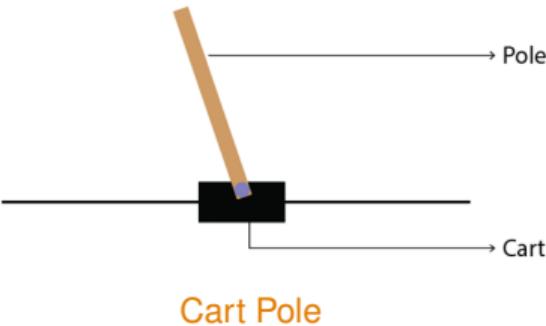


Cart Pole

from https://docs.pytorch.org/tutorials/intermediate/reinforcement_q_learning.html



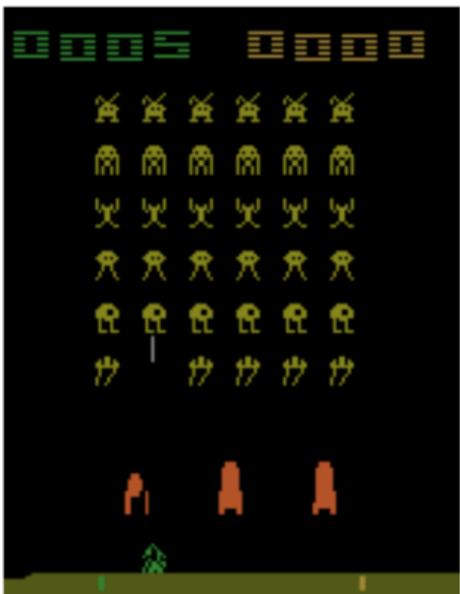
beginframe



| Element | Values set | | | | | | | | | | | | | | | | | | | | | | | |
|-------------------|---|------|-----|--|------|-------------|-----|-----|---|---------------|------|-----|---|---------------|------|-----|---|------------|------|-----|---|-----------------------|------|-----|
| Action Space | 0: Push cart to the left, 1: Push cart to the right | | | | | | | | | | | | | | | | | | | | | | | |
| Observation Space | <table border="1"><thead><tr><th>Num.</th><th>Observation</th><th>Min</th><th>Max</th></tr></thead><tbody><tr><td>0</td><td>Cart Position</td><td>-4.8</td><td>4.8</td></tr><tr><td>1</td><td>Cart Velocity</td><td>-Inf</td><td>Inf</td></tr><tr><td>2</td><td>Pole Angle</td><td>-24°</td><td>24°</td></tr><tr><td>3</td><td>Pole Angular Velocity</td><td>-Inf</td><td>Inf</td></tr></tbody></table> | | | | Num. | Observation | Min | Max | 0 | Cart Position | -4.8 | 4.8 | 1 | Cart Velocity | -Inf | Inf | 2 | Pole Angle | -24° | 24° | 3 | Pole Angular Velocity | -Inf | Inf |
| Num. | Observation | Min | Max | | | | | | | | | | | | | | | | | | | | | |
| 0 | Cart Position | -4.8 | 4.8 | | | | | | | | | | | | | | | | | | | | | |
| 1 | Cart Velocity | -Inf | Inf | | | | | | | | | | | | | | | | | | | | | |
| 2 | Pole Angle | -24° | 24° | | | | | | | | | | | | | | | | | | | | | |
| 3 | Pole Angular Velocity | -Inf | Inf | | | | | | | | | | | | | | | | | | | | | |
| Reward | +1 per step taken | | | | | | | | | | | | | | | | | | | | | | | |

beginframe





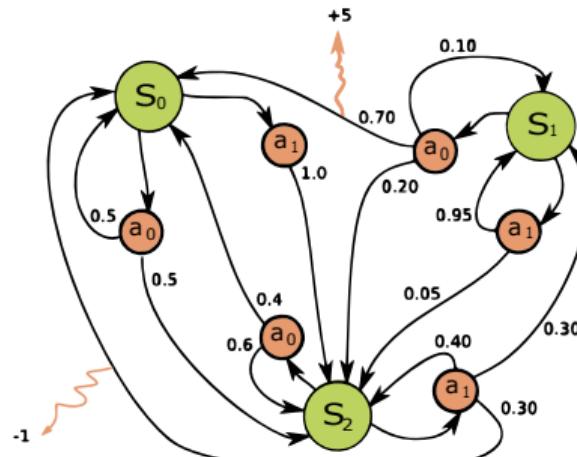
Space Invaders



| Element | Values set | | | |
|--------------------------|--|-------------------|--------------|----------------|
| Action Space | Value | Meaning | Value | Meaning |
| | 0 | NOOP | 1 | FIRE |
| | 2 | RIGHT | 3 | LEFT |
| | 4 | RIGHTFIRE | 5 | LEFTFIRE |
| Observation Space | Obs. Type | Dimensions | Range | |
| | ram | (128,) | [0-255] | |
| | rgb | (210, 160, 3) | [0-255] | |
| | grayscale | (210, 160) | [0-255] | |
| Reward | You gain points for destroying space invaders. The invaders in the back rows are worth more points. | | | |

RL modeled as Markov Decision Process (MDPs)

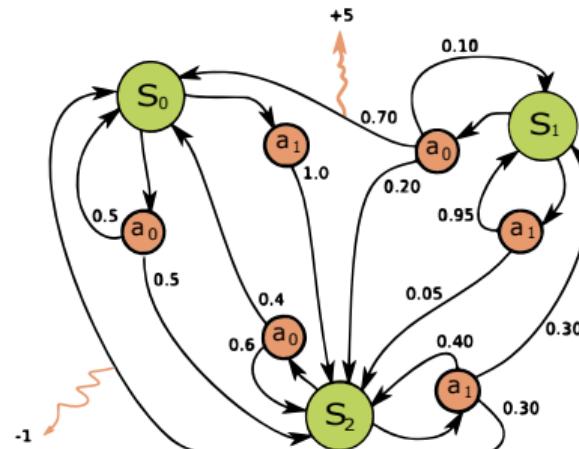
- Markov Decision Processes(MDPs) are the most common models for RL



Example of a Markov Decision Process, (author: waldoalvarez, licence Creative Commons)

RL modeled as Markov Decision Process (MDPs)

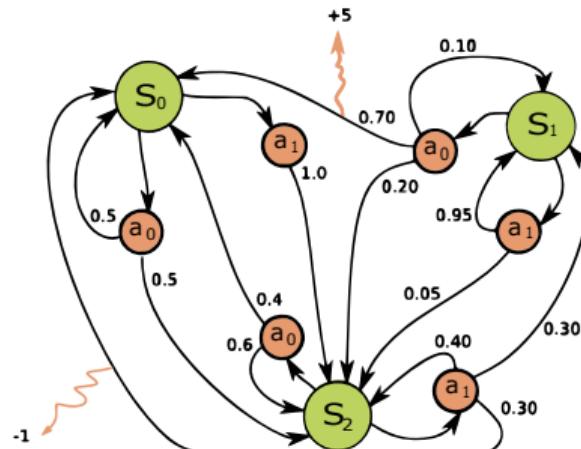
- Markov Decision Processes(MDPs) are the most common models for RL
- Markov assumptions:
 - Information state: sufficient statistic of history
 - State s_t is Markov if and only if:
 $p(s_{t+1}|s_t, a_t) = p(s_{t+1}|h_t, a_t)$ where History
 $h_t = (a_1, s_1, r_1, \dots, a_t, s_t, r_t)$
 - Future is independent of past given present



Example of a Markov Decision Process, (author: waldoalvarez, licence Creative Commons)

RL modeled as Markov Decision Process (MDPs)

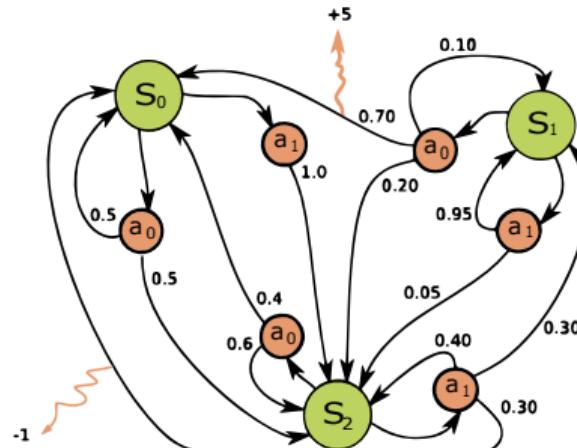
- Markov Decision Processes(MDPs) are the most common models for RL
- Markov assumptions:
 - Information state: sufficient statistic of history
 - State s_t is Markov if and only if:
 $p(s_{t+1}|s_t, a_t) = p(s_{t+1}|h_t, a_t)$ where History
 $h_t = (a_1, s_1, r_1, \dots, a_t, s_t, r_t)$
 - Future is independent of past given present
- behavior of an agent is defined by a **policy**



Example of a Markov Decision Process, (author: waldoalvarez, licence Creative Commons)

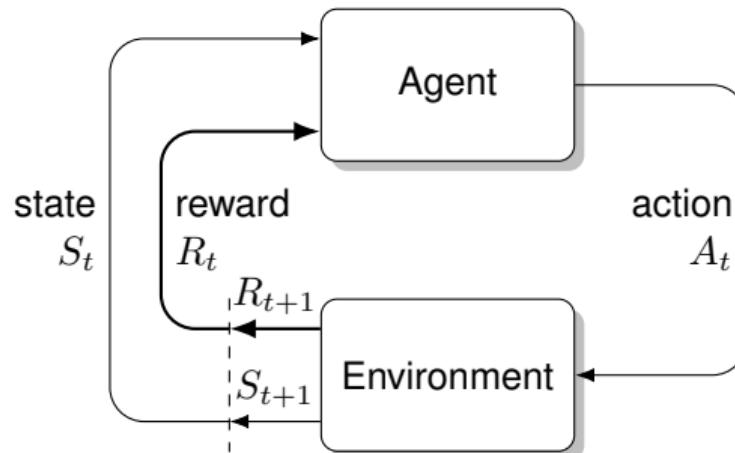
RL modeled as Markov Decision Process (MDPs)

- Markov Decision Processes(MDPs) are the most common models for RL
- Markov assumptions:
 - Information state: sufficient statistic of history
 - State s_t is Markov if and only if:
 $p(s_{t+1}|s_t, a_t) = p(s_{t+1}|h_t, a_t)$ where History
 $h_t = (a_1, s_1, r_1, \dots, a_t, s_t, r_t)$
 - Future is independent of past given present
- behavior of an agent is defined by a **policy**
- Goal of RL is to find the **optimal** policy



Example of a Markov Decision Process, (author: waldoalvarez, licence Creative Commons)

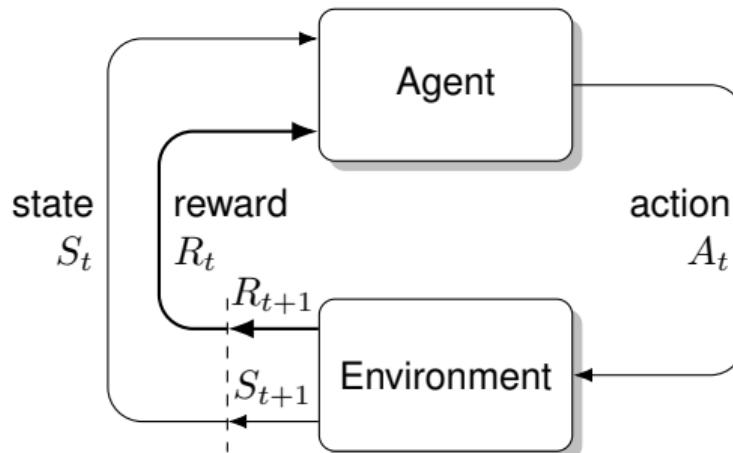
RL definitions



Agent-environment interaction loop in a RL system, from
[Sutton and Barto, 2018]



RL definitions

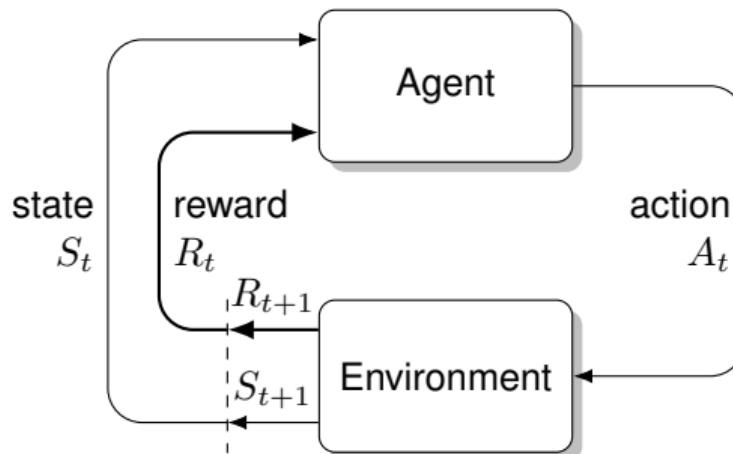


- **State (S_t):** Representation of the environment at any given time. Helps agent to make some decision.
- **Action (A_t):** Decision that an agent makes in a state to move in the environment

Agent-environment interaction loop in a RL system, from [Sutton and Barto, 2018]



RL definitions

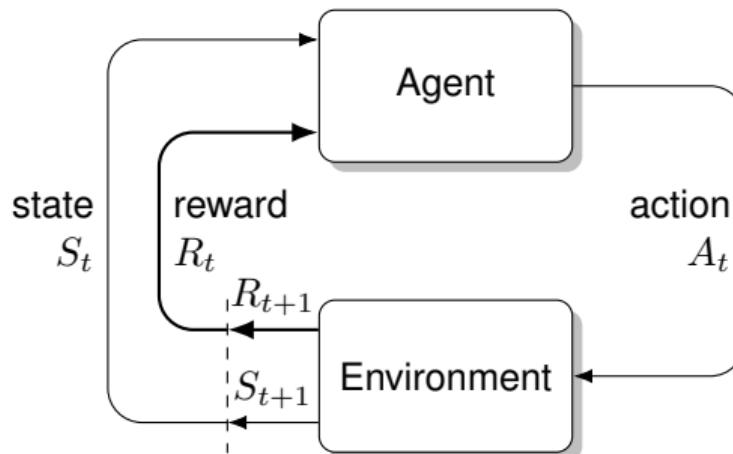


- **Reward (R_{t+1}):** Immediate feedback from the environment, scalar signal for taking some action.
- **Return (G):** Cumulative reward over an episode/terminal condition
$$G_t = R_t + R_{t+1} \dots + R_T$$

Agent-environment interaction loop in a RL system, from
[Sutton and Barto, 2018]



RL definitions

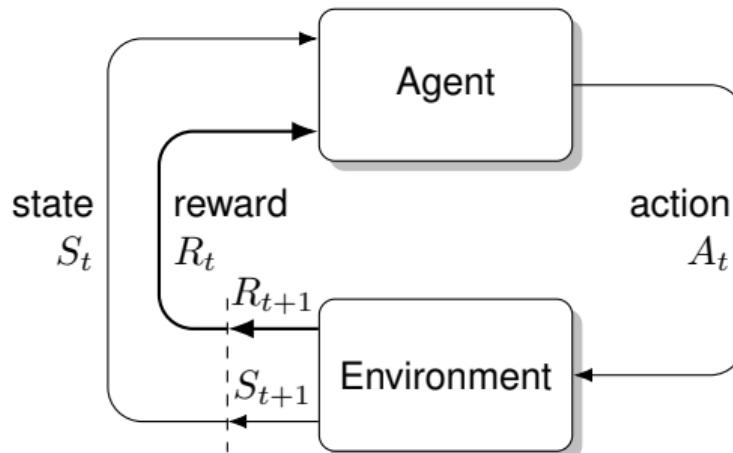


- **Policy (π)** determines how the agent chooses actions. $\pi : S \rightarrow A$
 - Deterministic policy: $\pi(s) = a$
 - Stochastic policy: $\pi(a|s) = P(a_t = a|s_t = s)$
- **Optimal Policy:** What agent has to learn ultimately. i.e. a policy that maximizes the return

Agent-environment interaction loop in a RL system, from [Sutton and Barto, 2018]



RL definitions



Agent-environment interaction loop in a RL system, from [Sutton and Barto, 2018]



MDP property

Theorem

For any Markov Decision Process:

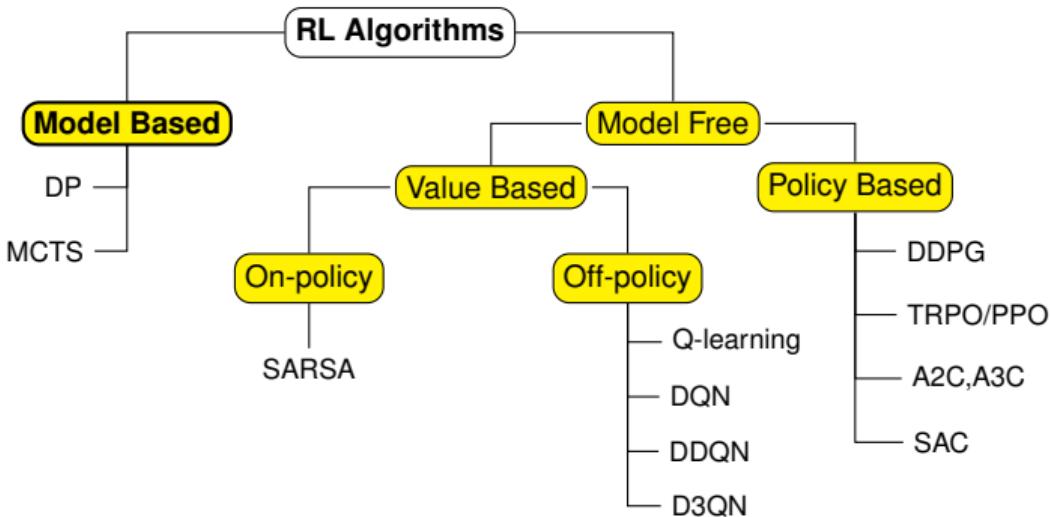
- There exists an Optimal Policy π^* , i.e., there exists a Policy π^* such that $V_\pi^*(s) \geq V_\pi(s)$ for all policies π and for all states $s \in S$
- All Optimal Policies achieve the Optimal Value Function, i.e. $V_\pi^*(s) = V^*(s)$ for all $s \in S$, for all Optimal Policies π^*
- All Optimal Policies achieve the Optimal Action-Value Function, i.e. $Q_\pi^*(s, a) = Q^*(s, a)$ for all $s \in S$, for all $a \in A$, for all Optimal Policies π^*

Cf. for example Prof. Ashwin Rao's course

https://web.stanford.edu/class/cme241/lecture_slides/OptimalPolicyExistence.pdf



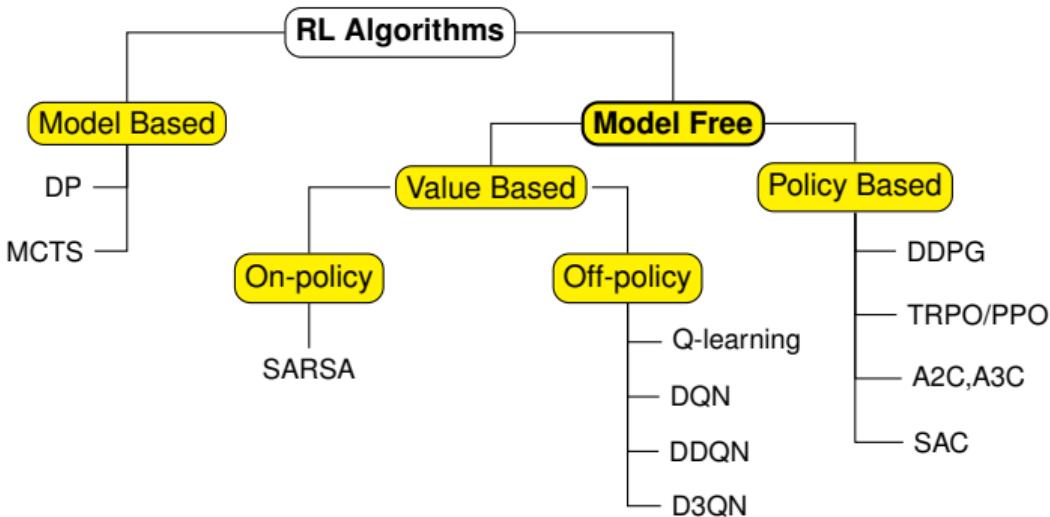
Taxonomy of algorithms



- **Model-based approaches:**
Dynamics and transition probabilities of the environment are known. i.e. state in which an action will lead with known reward



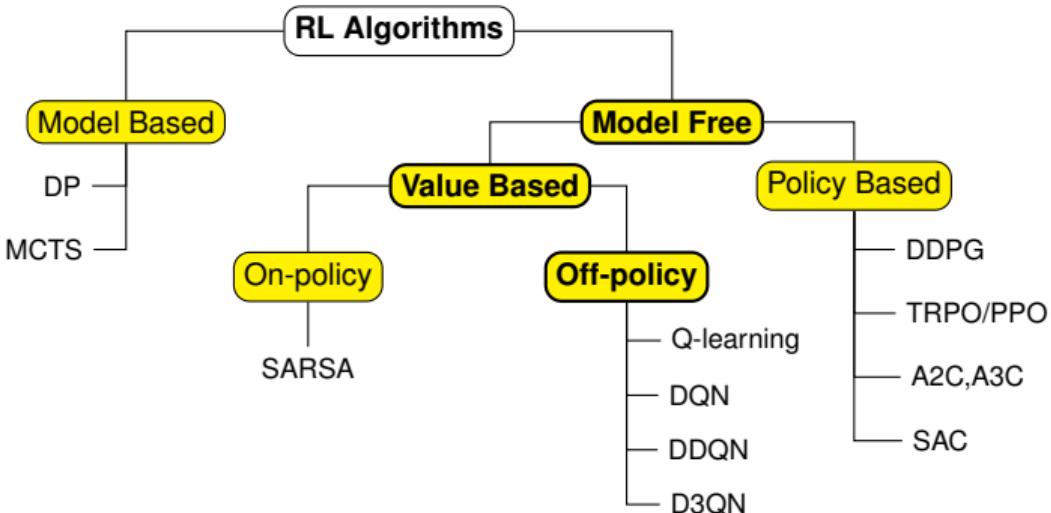
Taxonomy of algorithms



- **Model free approaches:** Most of the real world problems falls into this category. Internal dynamics are unknown. The only way to learn is through experience.



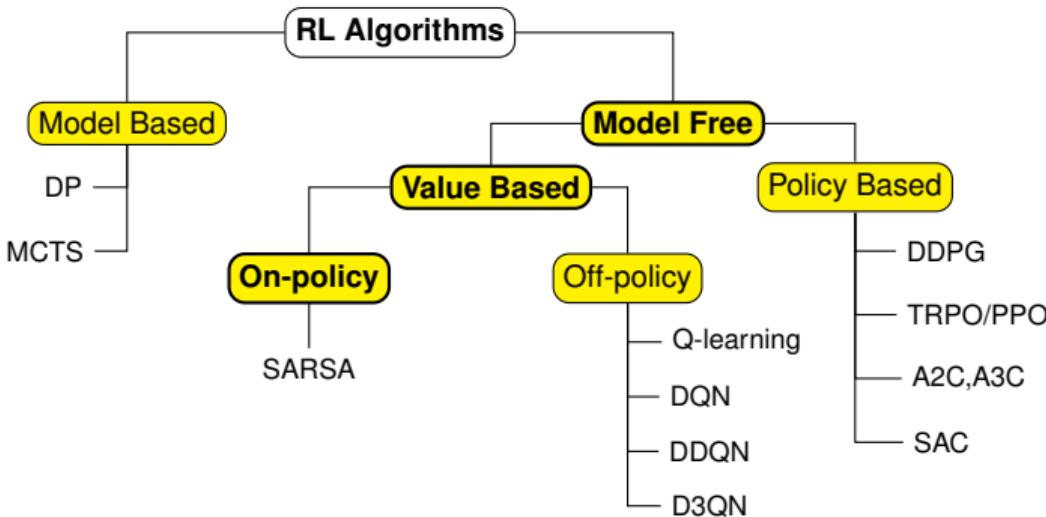
Taxonomy of algorithms



- **Model free approaches:** Most of the real world problems falls into this category. Internal dynamics are unknown. The only way to learn is through experience.
 - **Value Based:** Methods that learns the value function, which in turn derives an optimal policy
 - **Off-Policy:** Policy learned (*Target policy*) is different from policy used for exploring environment (*Behavior policy*).



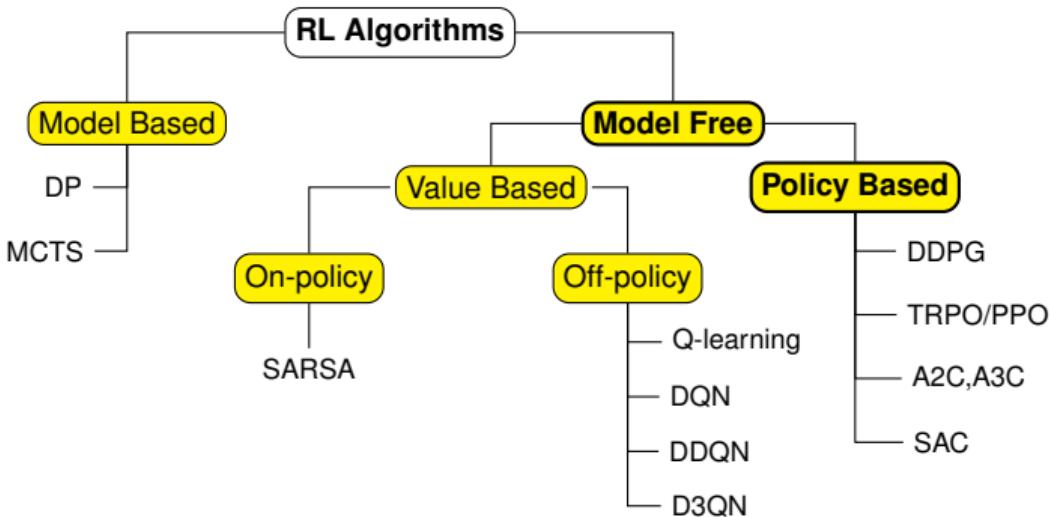
Taxonomy of algorithms



- **Model free approaches:** Most of the real world problems falls into this category. Internal dynamics are unknown. The only way to learn is through experience.
 - **Value Based:** Methods that learns the value function, which in turn derives an optimal policy
 - **On-Policy:** Target policy and behavior policy are same.



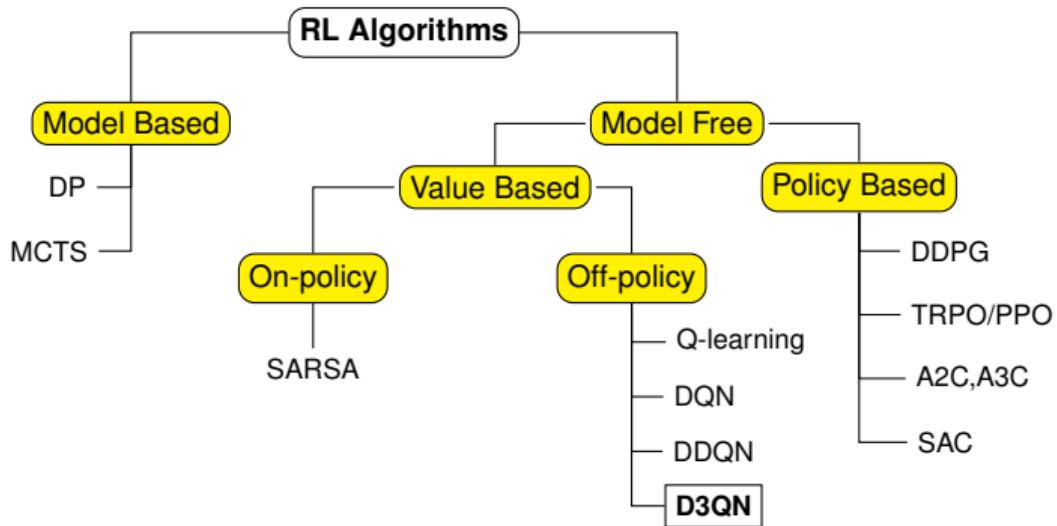
Taxonomy of algorithms



- **Policy Based:** Learns the optimal policy directly



Taxonomy of algorithms



Some fundamental equations in RL

- **Bellman's Equation:** A way/rule to recursively represent the Q-function or value function. It is an aggregation of two parts:

- **For value function:** Sum of immediate reward and return from next state

$$V_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma V_{\pi}(s_{t+1})]$$

- **For Q-function:**

- Immediate reward obtained from taking action a in state s to reach s_{t+1}

- Discounted return from next state s_{t+1} on executing action a_{t+1} & thereafter following policy π

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma Q_{\pi}(s_{t+1}, a_{t+1})]$$



Some fundamental equations in RL

- **Bellman's Equation:** A way/rule to recursively represent the Q-function or value function. It is an aggregation of two parts:

- **For value function:** Sum of immediate reward and return from next state

$$V_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma V_\pi(s_{t+1})]$$

- **For Q-function:**

- Immediate reward obtained from taking action a in state s to reach s_{t+1}

- Discounted return from next state s_{t+1} on executing action a_{t+1} & thereafter following policy π

$$Q_\pi(s, a) = \mathbb{E}_\pi[R_{t+1} + \gamma Q_\pi(s_{t+1}, a_{t+1})]$$

- **Temporal Difference learning (TD-learning):** method to learn the value function based on the difference between successive value estimates. As per the TD-learning rule the value function at each time step is updated as:

$$V_\pi(S_t) \leftarrow V_\pi(S_t) + \alpha[R_{t+1} + \gamma V_\pi(S_{t+1}) - V_\pi(S_t)]$$

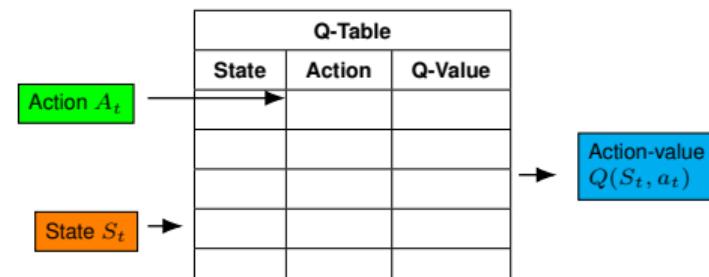
- **TD-target:** $R_{t+1} + \gamma V_\pi(S_{t+1})$
- **TD-error:** $R_{t+1} + \gamma V_\pi(S_{t+1}) - V_\pi(S_t)$



Q-learning

Q-learning [Watkins, 1989], [Watkins and Dayan, 1992]:

- Learning of the *Q* (or *action-value*) function
- **Model-free & off-policy** Temporal Difference (TD) control algorithm
- Tabular method. $Q \leftarrow S \times A$
- The optimal policy is given as:
$$\pi^* = \arg \max_a Q_\pi(s, a)$$



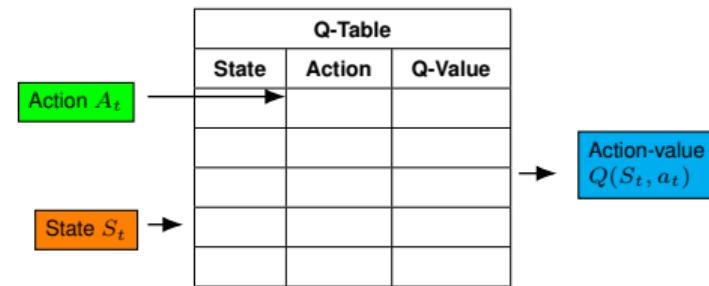
Q-table for *Q*-learning



Q-learning

- Q-values are recursively updated as the agent interacts with the environment
- ϵ -greedy policy:
 - **Exploration:** Random actions with probability ϵ .
 - **Exploitation:** as ϵ value decays agent follow the learned policy
- The Q-values are updated as per the TD-learning rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_t + \gamma * \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

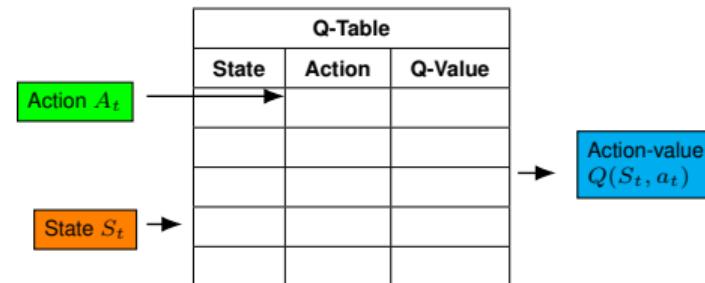


Q-table for *Q*-learning

Q-learning

- Q-values are recursively updated as the agent interacts with the environment
- ϵ -greedy policy:
 - **Exploration:** Random actions with probability ϵ .
 - **Exploitation:** as ϵ value decays agent follow the learned policy
- The Q-values are updated as per the TD-learning rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_t + \gamma * \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

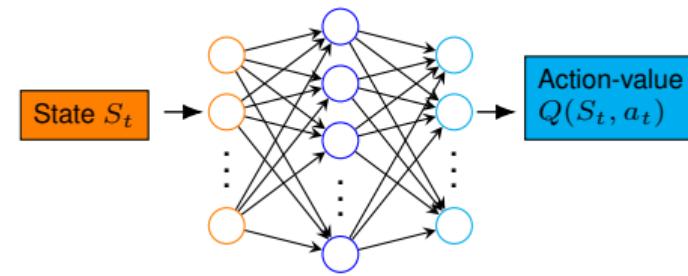


Q-table for *Q*-learning

Does not scale well if the state space is
or large: *the curse of dimensionality*

Deep Q-Network

- Deep Q-Network (DQN) [Mnih et al., 2015] makes use of Deep Neural Networks (NNs) to approximate the optimal Q-function
- DQN learns to output the optimal Q-values by minimizing the loss $Q_{predicted} - Q_{actual}$ i.e. TD error $R_{t+1} + \gamma V_\pi(S_{t+1}) - V_\pi(S_t)$.
- It uses two networks:
 - **Policy Network:** To generate the Q-values for actions taken in the current state
 - **Target Network:** To create the Q-values for the next stateTarget network is updated after some epochs with the Policy network



Deep Q-Network (DQN)

Double Deep Q-Networks

- DQN suffers from overestimation bias & moving target problems.
- DDQN overcomes them by using a different method of **TD target formation**
- Decouples action selection from target Q-value generation. It uses 2 networks similar to DQN.
 - uses Policy network to select the best action for next state (θ)
 - Q-value for the selected action for next state is generated using the target network. (θ')
- In the form of Bellman's equation

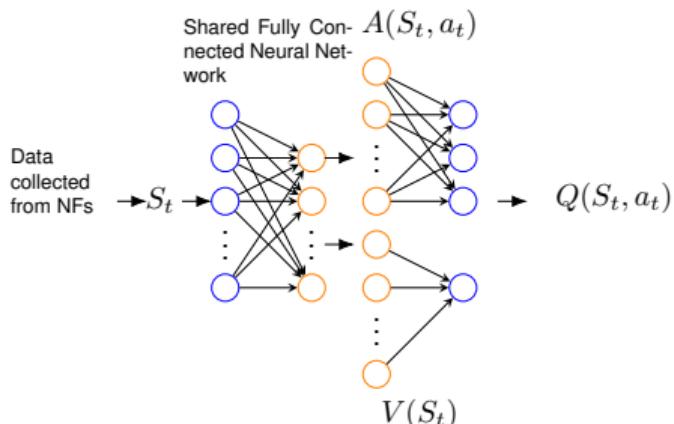
$$Q(s, a; \theta) = r + \gamma \max_{a'} Q(s', a'; \theta')$$



D3QN

Dueling Double Deep Q-networks (D3QN) [Wang et al., 2016] is an architectural enhancement to improve learning efficiency of Q-function:

- It combines the functionality of DQN & DDQN.
- D3QN can learn which states are valuable without considering the effect of each action.
- It decomposes the Q-function as the sum of:
 - **Value stream $V(s)$:** Estimate of how good it is to be in a state. $V(s)$ is estimated by one separate stream
 - **Advantage stream $A(s,a)$:** How much extra reward an action can result into. $A(s,a)$ is calculated by another stream

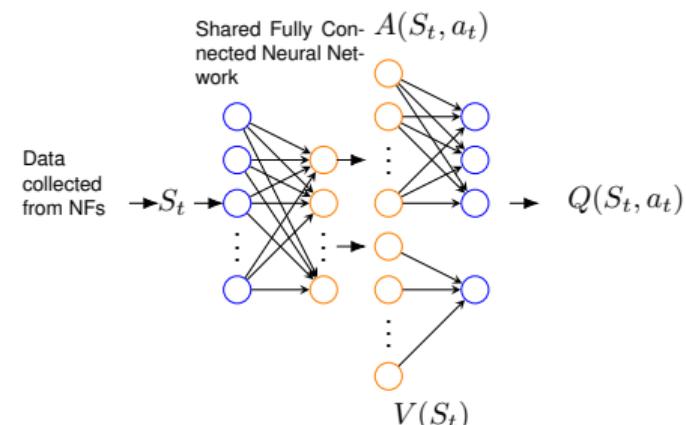


Overview of D3QN

D3QN

D3QN [Wang et al., 2016] is an architectural enhancement to improve learning efficiency of Q-function:

- It combines the functionality of DQN & DDQN.
- D3QN can learn which states are valuable without considering the effect of each action.
- $V(s)$ and $A(s,a)$ are combined through an aggregation layer. $Q(s, a) = V(s) + A(s, a)$
- $Q(S, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + (A(s, a; \theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum A(s, a'; \theta, \alpha))$

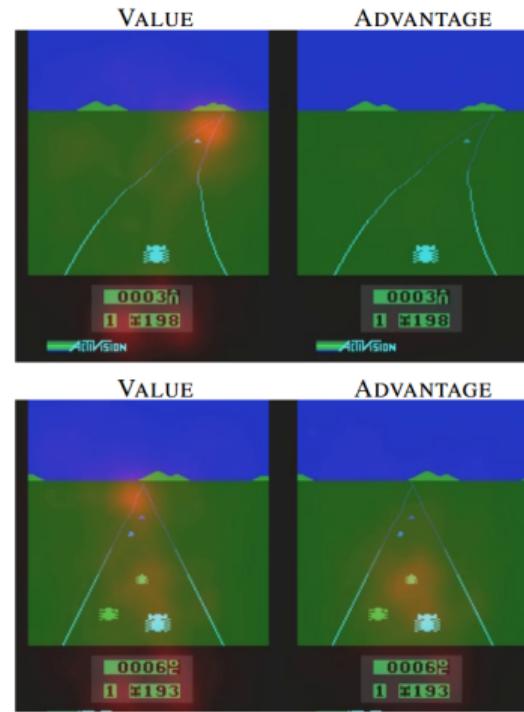


Overview of D3QN



D3QN: Advantage and Value streams behaviors

The advantage stream learns to pay attention only when there are cars immediately in front, so as to avoid collisions. (from [Wang et al., 2016])



Some further details

Prioritized Experience Replay (PER):

- **Experience Replay Buffer** from [Lin, 1992]: stores the experiences (s, a, r, s') collected by agent's interaction with the environment.
 - Common approach in Deep Q-networks.
 - Data is sampled randomly or uniformly from replay buffer to train the DNN
 - Helps in breaking in temporal correlation present in the sequential data
 - Introduces sample efficiency: sampling data from buffer ensures to have independent and identically distributed (i.i.d.) data.
- PER samples the experiences by prioritizing the samples with higher TD-error. By focusing more on challenging situation it helps with:
 - Improving sample efficiency
 - Faster convergence



Some further details

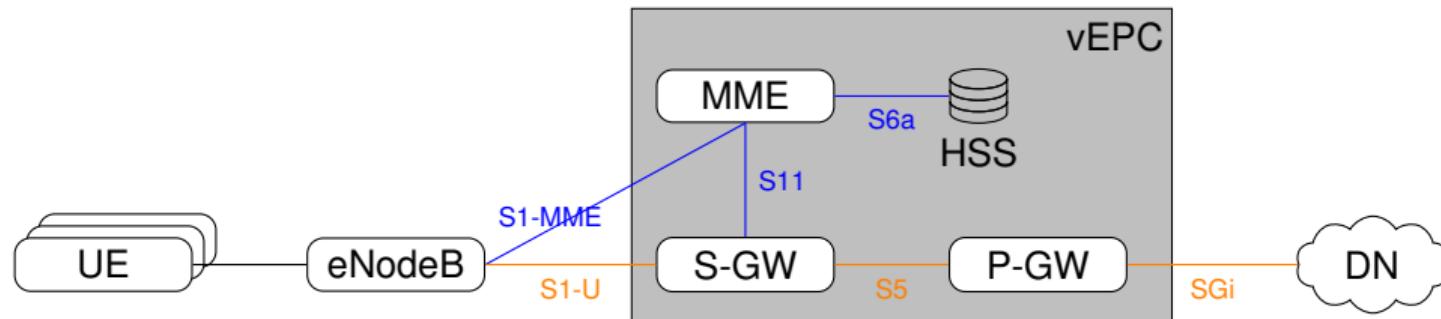
Action Masking (AM) from [Huang and Ontañón, 2022] :

- Additional extension used in our experiments
- Some actions that are not allowed are masked during the exploration phase
- Example: turning off the last instance of a Network Function (NF) when one requirement is to always have at least one instance running



Scalability problem

- The scalability problem for NFs is the scaling of cloud resources allocated to NFs according to the workload to guarantee the Quality of Service (QoS).
- This example focus on the horizontal scalability of virtualized Evolved Packet Core (vEPC) NFs
- Objective is to automatically scale the platform depending on the user traffic

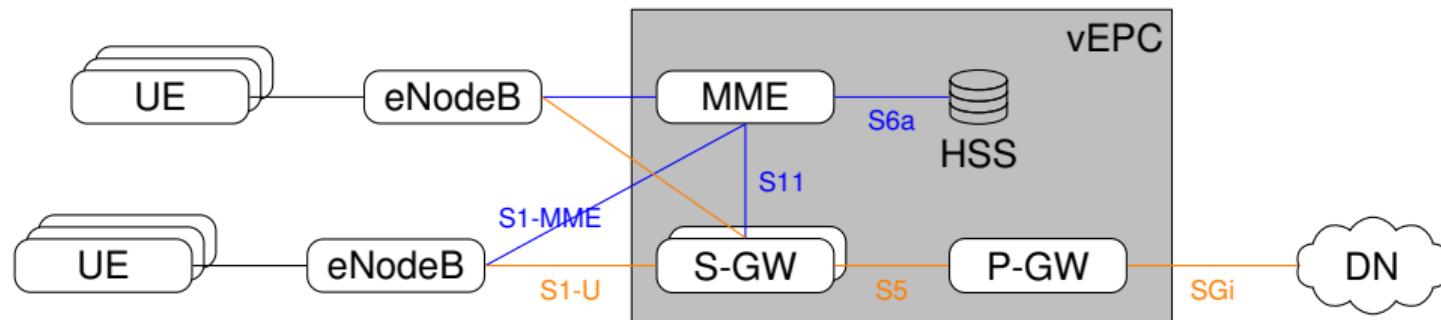


Problem description for LTE NFs



Scalability problem

- The scalability problem for NFs is the scaling of cloud resources allocated to NFs according to the workload to guarantee the QoS.
- This example focus on the horizontal scalability of vEPC NFs
- Objective is to automatically scale the platform depending on the user traffic

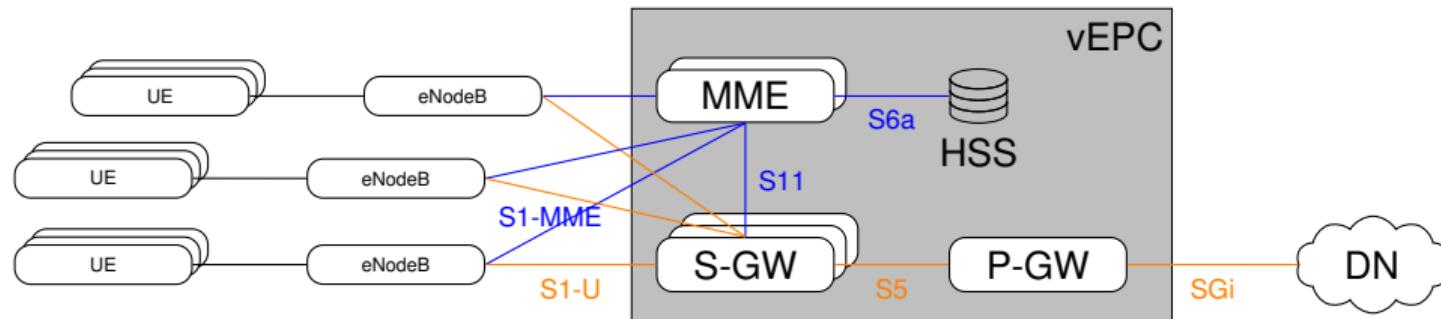


Problem description for LTE NFs



Scalability problem

- The scalability problem for NFs is the scaling of cloud resources allocated to NFs according to the workload to guarantee the QoS.
- This example focus on the horizontal scalability of vEPC NFs
- Objective is to automatically scale the platform depending on the user traffic



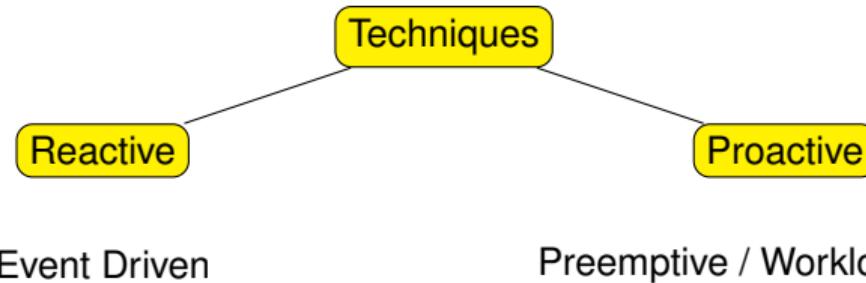
Problem description for LTE NFs



A short introduction of the problem statement

Autoscaling autonomously & dynamically provision (or deprovision) a set of resources in a way that

- caters fluctuant application workloads
- optimize the overall cost
- satisfies application SLA or SLO



Reactive vs Proactive autoscaling

Reactive: where system scales in real-time in response to the workload

- Threshold or event based, where resources are scaled as per pre-defined rules, e.g. scale out when CPU utilization reaches 80%, scale in when it reaches 20%
- Suitable for smooth and gradual workloads
- Pros:
 - Straightforward to implement
 - Scaling on-demand can save some cost & ensure availability
- Cons:
 - Resource flapping: scaling oscillations due to rapidly changing workload
 - longer provisioning time of resources can degrade SLA



Reactive vs Proactive Scaling

Proactive: when a system adjusts itself preemptively to account for upcoming workload variations.

- Generally involves use of control theory, queuing theory, predictive models or machine learning techniques. For example:
- Suitable for systems where SLA degradation and downtime are not acceptable.
 - Time series: forecasts the workload or resource consumption to scale ahead of time.
 - **Reinforcement Learning:** an agent learns a scaling policy by trial and error so that the policy will help achieve optimization objective
- Pros:
 - Scaling ahead of demand ensures application availability and QoS
 - Avoid taking too many actions when the load fluctuates a lot
- Cons:
 - More complex to design
 - Might be difficult to get approval from operational teams



Motivation for using RL for this use case

- Autoscaling can be referred to a classic automation control problem
- Commonly abstracted as MAPE control loop, which repeats itself over time [[[Kephart and Chess, 2003](#)]]
 - Monitoring: monitor the performance indicators at regular intervals
 - Analysis: whether it is necessary to scale based on monitored KPIs
 - Planning: how many resources to provision/deprovision
 - Execution: execute the plan using cloud provider APIs
- **RL addresses optimal control problems**
 - Capability of learning effect of action over long-term
 - Optimization is a sequential decision making process.



Reward function

| Metric | Definition |
|--------|---|
| U | Number of User Equipments (UEs) connected |
| M | Memory usage, in MB |
| D | Number of dropped sessions |

- Maximum reward value is 1
- Encourages the NFs to use resources optimally around 70%
- Resource usage above 80% or crashes are penalized with lowest reward value –1.

$$r = \begin{cases} 1 - (0.7 - \max(M, U) - D) & \text{if } M, U \in [0, 80] \\ \max(-\max(M, U) - 10 * D, -1) & \text{otherwise.} \end{cases}$$



Lessons learned

RL is complex:

- Not your typical network or software engineer skill
- Very active area of research



Lessons learned

RL is complex:

- Not your typical network or software engineer skill
- Very active area of research

Automation of Network Core scaling:

- Development is required
- Probably better to use Infrastructure as Code framework to be IaaS-independent



Lessons learned

RL is complex:

- Not your typical network or software engineer skill
- Very active area of research

Automation of Network Core scaling:

- Development is required
- Probably better to use Infrastructure as Code framework to be IaaS-independent

Load generation is always tricky:

- Commercial products exist
- Lack of open source tools to generate traffic in a consistent way for a long time
- Traffic needs to be balanced across all instances



Resources on RL

- Reinforcement Learning: An Introduction (second edition) by Richard S. Sutton and Andrew G. Barto. Licensed under the Creative Commons Attribution-NonCommercial-NoDerivs 2.0 Generic License
- Stanford CS234: Reinforcement Learning | Spring 2024 the complete CS234: Reinforcement Learning course of Prof. Emma Brunskill at Stanford captured on video
- David Silver's course it includes a link to the video captures
- Reinforcement Learning Specialization by Martha White and Adam White at the University of Alberta, available on Coursera.

Bibliography

- Berner, C., Brockman, G., Chan, B., Cheung, V., Dębiak, P., Dennison, C., Farhi, D., Fischer, Q., Hashme, S., Hesse, C., et al. (2019). Dota 2 with large scale deep reinforcement learning. *arXiv preprint arXiv:1912.06680*.
- Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., and Meger, D. (2018). Deep reinforcement learning that matters. In *Proceedings of the AAAI conference on artificial intelligence*, volume 32.
- Huang, S. and Ontañón, S. (2022). A closer look at invalid action masking in policy gradient algorithms. *The International FLAIRS Conference Proceedings*, 35.
- Kephart, J. O. and Chess, D. M. (2003). The vision of autonomic computing. *Computer*, 36(1):41–50.
- Lin, L.-J. (1992). Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine learning*, 8:293–321.
- Mitchell, T. (1997). Learning, machine. *McGraw Hill*.
- Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing Atari with Deep Reinforcement Learning.
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., Petersen, S., Beattie, C., Sadik, A., Antonoglou, I., King, H., Kumaran, D., Wierstra, D., Legg, S., and Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533.
- Qiu, H., Banerjee, S. S., Jha, S., Kalbarczyk, Z. T., and Iyer, R. K. (2020). FIRM: An intelligent fine-grained resource management framework for SLO-Oriented microservices. In *14th USENIX Symposium on Operating Systems Design and Implementation (OSDI 20)*, pages 805–825. USENIX Association.

Bibliography (cont.)

- Qiu, H., Mao, W., Patke, A., Wang, C., Franke, H., Kalbarczyk, Z. T., Başar, T., and Iyer, R. K. (2022). SIMPPO: A scalable and incremental online learning framework for serverless resource management. In *Proceedings of the 13th Symposium on Cloud Computing*, SoCC '22, page 306–322, New York, NY, USA. Association for Computing Machinery.
- Sutton, R. S. and Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT Press.
- Vinyals, O., Babuschkin, I., Czarnecki, W. M., Mathieu, M., Dudzik, A., Chung, J., Choi, D. H., Powell, R., Ewalds, T., Georgiev, P., et al. (2019). Grandmaster level in Starcraft II using multi-agent reinforcement learning. *Nature*, 575(7782):350–354.
- Wang, Z., Schaul, T., Hessel, M., Hasselt, H., Lanctot, M., and Freitas, N. (2016). Dueling Network Architectures for Deep Reinforcement Learning. In *Proceedings of The 33rd International Conference on Machine Learning*, pages 1995–2003. PMLR.
- Watkins, C. J. C. H. (1989). Learning from delayed rewards. *PhD dissertation*.
- Watkins, C. J. C. H. and Dayan, P. (1992). Q-learning. *Machine Learning*, 8(3–4):279–292.